

IGV: Incremental Graph Versioning for Efficient Retrieval-Augmented Generation over Evolving Corpora

Anonymous Author(s)

ABSTRACT

Graph-based Retrieval-Augmented Generation (GraphRAG) enhances large language models (LLMs) by structuring document corpora as knowledge graphs, enabling multi-hop reasoning and global question answering. However, existing GraphRAG systems require expensive full-graph rebuilds when new documents arrive, making them impractical for dynamic, evolving corpora. We present IGV (Incremental Graph Versioner), a framework that supports efficient incremental updates to GraphRAG indexes through three components: (1) embedding-based semantic entity deduplication, (2) embedding-based graph relinking, and (3) incremental community repartitioning with community summary generation for global QA retrieval. Experiments on five multi-hop QA datasets with LLM-based entity extraction (Qwen2.5-7B-Instruct) across 2,000 passages and 120 evaluation questions per dataset demonstrate that IGV achieves 10 $\times$ faster updates than full rebuild while maintaining factoid QA quality. Critically, IGV’s community summary retrieval improves global QA F1 by 12.6–93.8% over naive incremental baselines on 3/5 datasets, and even surpasses VectorRAG on streaming corpora. All methods share cached LLM entity extractions, ensuring fair ablation comparisons.

KEYWORDS

GraphRAG, Incremental Indexing, Knowledge Graph, Retrieval-Augmented Generation

1 INTRODUCTION

Retrieval-Augmented Generation (RAG) mitigates LLM hallucinations by grounding responses in retrieved external knowledge. While traditional vector RAG excels at factoid lookups, it struggles with global questions requiring cross-document reasoning—questions such as “What are the main themes across this document collection?” or “How are these entities related?” GraphRAG addresses this limitation by constructing knowledge graphs from text corpora, enabling structured multi-hop retrieval and community-level summarization [7, 10]. By organizing entities and their relationships into a graph index, GraphRAG systems can answer queries that require synthesizing information across multiple documents, which flat vector retrieval cannot handle effectively.

A critical limitation of current GraphRAG systems is their reliance on *full rebuilds* for corpus updates. Microsoft GraphRAG [7] requires re-running entity extraction, community detection, and summary generation over the entire corpus when any new document arrives. For a 5,000-passages corpus, this costs over 4.5 hours of LLM time and significant

API expenditure. LightRAG [10] introduced a set-union incremental update mechanism that avoids full rebuilds, but its naive merging approach creates three problems: (1) duplicate entities—the same real-world entity extracted from different documents with slightly different surface forms creates redundant nodes; (2) disconnected graph components—new entities are added without connecting them to the existing graph structure; and (3) stale communities—community assignments computed on the initial corpus become outdated as new entities arrive, degrading global query performance.

These problems compound over time: with each incremental batch, the graph becomes increasingly bloated with duplicates, fragmented across disconnected components, and organized into communities that no longer reflect the current graph structure. This degradation directly impacts retrieval quality, and the cost of full rebuild makes GraphRAG impractical for streaming applications such as news monitoring, enterprise knowledge base maintenance, and scientific literature tracking.

We propose IGV (Incremental Graph Versioner), which addresses these problems through three components: semantic entity deduplication using sentence-transformer embeddings, embedding-based graph relinking that connects semantically similar entities, and incremental community repartitioning with LLM-generated community summaries for global QA retrieval. The key insight is that community summaries provide a thematic index for global questions, which factoid retrieval cannot serve effectively. Unlike prior work that treats community detection as an offline preprocessing step, IGV maintains communities incrementally and leverages them at query time through community summary retrieval.

Our contributions are as follows:

- We formulate the incremental GraphRAG maintenance problem and identify three challenges: entity duplication, graph fragmentation, and community staleness.
- We propose IGV, combining embedding-based deduplication, semantic relinking, and incremental community repartitioning with community summary generation.
- We introduce a fair ablation methodology using cached LLM entity extractions, eliminating the non-determinism that invalidated prior comparisons.
- We evaluate on five datasets with two QA modes (factoid and global), demonstrating 10 $\times$ efficiency and 12.6–93.8% global QA improvement on 3/5 datasets.

2 RELATED WORK

GraphRAG Systems. Edge et al. [7] proposed GraphRAG, which builds entity knowledge graphs via LLM extraction,

applies Leiden community detection [36], and generates community summaries for global query answering. Their system processes documents in two phases: offline indexing and on-line query. While effective for static corpora, the system requires complete re-indexing when new documents arrive. LightRAG [10] simplified this with a dual-level retrieval paradigm and introduced basic incremental updates via set union, but lacks deduplication, relinking, and incremental community maintenance. PathRAG [6] introduced flow-based path pruning to reduce retrieval redundancy. HippoRAG [15] used Personalized PageRank for neurobiologically inspired retrieval. GNN-RAG [26] combined graph neural networks with LLMs for KGQA. GRAG [14] addressed networked document retrieval. ReGraphRAG [16] reorganized fragmented knowledge graphs. GFM-RAG [24] proposed a graph foundation model for RAG. LEGO-GraphRAG [5] modularized the GraphRAG pipeline. Comprehensive surveys [11, 28] provide systematic overviews of the GraphRAG landscape. Think-on-Graph [32] and its successor [25] proposed interactive KG reasoning with LLMs. Simple is Effective [19] analyzed the roles of graphs and LLMs in KG-based RAG. G-Retriever [12] enabled chatting with textual graphs.

RAG Fundamentals and Applications. Lewis et al. [18] introduced RAG for knowledge-intensive NLP tasks, establishing the retrieve-then-generate paradigm. KG-guided RAG [50] used KG semantics for query disambiguation. Paths-over-Graph [33] extracted reasoning paths from KGs. Domain-specific applications include medical [22, 41, 48], financial [31], customer service [43], recommendation [37], and multimodal [4, 17, 21] settings. KG alignment [34] and personalization [29] have also been explored. LLM-based KG construction [13, 46] and hallucination reduction [1] are active research areas. Graph-of-Thoughts [2], MindMap [40], and graph neural prompting [35] bridge graph structures with LLM reasoning. Generate-on-Graph [44] treats LLMs as both agent and KG.

Incremental and Continual Learning. GORAG [38] proposed online graph updates for social media classification, addressing dynamic graph construction but without community-level reasoning. KAG [20] introduced logical-form-guided reasoning with schema constraints. Practical GraphRAG [27] reduced construction cost via dependency parsing. Continual graph learning [45] and dynamic GNNs [49] address evolving graph structures but focus on representation learning rather than retrieval. Continual learning for LLMs [42] surveys knowledge update strategies. Incremental learning approaches include exemplar replay [47], sub-graph distillation [8], and high-order experience replay [39]. Temporal KGQA methods [9, 30] handle time-aware reasoning. GNN explainability [23] is studied for model transparency. Community detection algorithms [3, 36] provide the algorithmic foundation for our incremental repartitioning. Louvain [3] optimizes modularity through local node reassignment, while Leiden [36] guarantees well-connected communities. Both operate on the full graph, making incremental variants essential for dynamic corpora.

3 METHODOLOGY

3.1 Problem Formulation

Given a knowledge graph $G = (V, E)$ constructed from an initial document corpus $\mathcal{D}_0$, and a stream of new document batches $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_T$ arriving over time, the goal is to update G to incorporate new entities and relationships from each batch $\mathcal{D}_t$ while satisfying four objectives:

- (1) **Efficiency:** Minimizing LLM calls per update, the dominant cost in GraphRAG.
- (2) **Consistency:** Avoiding duplicate entities that fragment the graph.
- (3) **Community quality:** Maintaining community structure that reflects the current graph topology.
- (4) **Retrieval utility:** Preserving factoid QA quality while improving global QA through community-level reasoning.

The full rebuild approach processes all documents $\bigcup_{i=0}^t \mathcal{D}_i$ from scratch at each step, with cost $O(\sum_{i=0}^t |\mathcal{D}_i|)$ LLM calls. The naive incremental approach (LightRAG) processes only the new batch $\mathcal{D}_t$, with cost $O(|\mathcal{D}_t|)$, but sacrifices graph consistency and community quality. IGV achieves the same $O(|\mathcal{D}_t|)$ LLM cost while maintaining consistency through deduplication, connectivity through relinking, and community quality through incremental repartitioning.

3.2 IGV Architecture

IGV processes each incremental batch $\mathcal{D}_t$ through four sequential stages, as formalized in Algorithm 1.

Stage 1: LLM-Based Entity Extraction. For each new document $d \in \mathcal{D}_t$, Qwen2.5-7B-Instruct extracts named entities and their relationships as triplets (s, r, o) , where s and o are entity mentions and r is the relationship description. The extraction prompt instructs the LLM to identify people, organizations, locations, concepts, and other named entities, returning results as a JSON list. Each entity is augmented with a description derived from its surrounding context. A document node (DOC- d) is created for each document, connecting all entities extracted from it. A critical design choice is that entity extraction results are *cached* and shared across all methods, ensuring that ablation comparisons reflect only graph construction differences, not LLM non-determinism. This is essential because LLM-based extraction is non-deterministic: without caching, the same document may yield different entity sets across runs, introducing noise that can exceed the effect of the components being ablated.

Stage 2: Semantic Deduplication. After extraction, new entities are compared against existing graph nodes to identify potential duplicates. We encode all entity names and descriptions using a sentence-transformer model (all-MiniLM-L6-v2, $d = 384$ dimensions) with L2-normalized embeddings. Given a new entity e_{new} with embedding $\mathbf{e}_{\text{new}}$ and existing entities with embeddings $\{\mathbf{e}_1, \dots, \mathbf{e}_N\}$, we compute:

$$\text{sim}(e_{\text{new}}, e_i) = \mathbf{e}_{\text{new}} \cdot \mathbf{e}_i \quad (1)$$

where the dot product equals cosine similarity due to L2 normalization. If $\text{sim}(e_{\text{new}}, e_i) \geq \tau$ (threshold $\tau = 0.95$), the entities are merged: the new entity’s relationships are redirected to the existing entity, descriptions are concatenated (keeping the longer one as primary), and source document IDs are accumulated. Document nodes (DOC_ prefix) are excluded from deduplication to prevent cross-document merging, since document nodes represent unique documents and should never be merged. This conservative threshold avoids over-merging distinct entities that happen to share similar descriptions.

Stage 3: Embedding-Based Graph Relinking. After deduplication, new entities may be disconnected from the existing graph if their source documents do not share entities with previously indexed documents. Unlike BFS-based relinking that only connects structurally proximate nodes (which fails when new entities are in isolated components), we use embedding-based relinking: for each new entity, we compute embedding similarity against all existing non-DOC_ entities and add edges to the top-3 most similar entities with cosine similarity ≥ 0.75 . This creates semantically meaningful cross-document connections that enable multi-hop reasoning paths across document boundaries.

Stage 4: Incremental Community Repartitioning. Rather than recomputing communities for the entire graph, IGV identifies the *affected region*—the set of nodes whose community membership may have changed. This region consists of all new nodes plus their 2-hop neighbors in the existing graph. We extract the subgraph induced by this region and re-apply the Louvain algorithm [3] only on this subgraph. The new sub-communities are then merged back into the global partition: existing community assignments for unaffected nodes are preserved, while affected nodes are reassigned to their new communities. After repartitioning, IGV generates LLM summaries for each community by providing the entity names and sample passages to Qwen2.5-7B-Instruct. These summaries serve as a thematic index for global QA retrieval.

3.3 Retrieval

IGV supports two retrieval modes, selected based on question type.

Factoid Retrieval uses dual-level graph retrieval: (1) *low-level* entity matching via embedding similarity—the query is encoded and compared against all entity embeddings, top-30 matching entities are selected, and passages containing these entities receive scores weighted by similarity and inverse degree; (2) *1-hop expansion*—neighbors of matched entities also receive passage scores at a lower weight ($0.5\times$); and (3) *high-level* relation edge matching—edges whose descriptions overlap with query terms boost passages containing either endpoint. This serves questions like “Who directed the film *Ed Wood*?”

Global Retrieval extends factoid retrieval with a community summary bonus. Community summaries are embedded and compared to the query via cosine similarity. Passages containing entities from the top-5 most similar communities

Algorithm 1 IGV Incremental Update

Require: New batch $\mathcal{D}_t$, cached entities $\mathcal{C}$, graph G , communities $\mathcal{C}_{t-1}$
Ensure: Updated graph G' , communities $\mathcal{C}_t$, summaries $\mathcal{S}_t$

- 1: **Stage 1:** Extract entities from $\mathcal{D}_t$ via LLM (or load from $\mathcal{C}$)
- 2: **Stage 2:** Encode new entities, compute sim vs existing (Eq. 1)
- 3: Merge entities with $\text{sim} \geq 0.95$ (exclude DOC_ nodes)
- 4: Add non-merged entities and triplets to G
- 5: **Stage 3:** For each new entity, find top-3 similar existing ($\text{sim} \geq 0.75$)
- 6: Add relink edges to G
- 7: **Stage 4:** Identify affected region S_t (new nodes + 2-hop neighbors)
- 8: $\mathcal{C}_t \leftarrow \text{Louvain}(G[S_t]) \cup (\mathcal{C}_{t-1} \setminus S_t)$
- 9: Generate LLM summaries $\mathcal{S}_t$ for each community in $\mathcal{C}_t$
- 10: **return** $G, \mathcal{C}_t, \mathcal{S}_t$

Table 1: Complexity per incremental batch (m : new entities, N : total, $|S|$: affected subgraph).

Operation	Full Rebuild	IGV
Entity extraction	$O(N + m)$	$O(m)$
Deduplication	$O((N + m)^2)$	$O(m \cdot N)$
Community detection	$O(N \log N)$	$O(S \log S)$
Community summaries	$O(N)$	$O(S)$
Total LLM calls	$O(N + m)$	$O(m)$

receive a bonus weight ($0.3\times$ community similarity score) *added on top of* the dual-level scores. The bonus is deliberately small to avoid degrading retrieval quality—communities provide a re-ranking signal, not a filter. This design ensures that global retrieval never performs worse than factoid retrieval in expectation, while providing additional context for abstract questions. This serves questions like “What are the main themes in this corpus?”

3.4 Complexity Analysis

Table 1 compares the computational complexity of IGV against full rebuild. When $m \ll N$ (the typical streaming case), IGV’s cost is dominated by the deduplication step $O(m \cdot N)$, which is significantly cheaper than the rebuild cost $O(N + m)$ for entity extraction alone. The community repartition cost $O(|S| \log |S|)$ is negligible since $|S| \approx m \cdot d^h \ll N$. Community summary generation costs $O(|S|)$ LLM calls, affecting only new communities rather than all communities.

4 EXPERIMENTS

4.1 Setup

Datasets. We evaluate on five multi-hop QA benchmarks: HotpotQA (Wikipedia-based multi-hop reasoning, 2,000 passages), 2WikiMultiHopQA (entity-centric multi-hop, 2,260 passages), MuSiQue (compositional multi-hop with 2–4 hops,

2,000 passages), NarrativeQA (long-form narrative understanding, 2,000 passages), and StreamingQA (temporally ordered passages simulating streaming updates, 2,000 passages). Each dataset is split into 70% base corpus (initial index) and three incremental batches of 10% each, simulating a streaming document arrival pattern.

LLM and Hardware. Qwen2.5-7B-Instruct for entity extraction, answer generation, and community summaries, running on a single NVIDIA RTX 5090 GPU (32GB VRAM) with bfloat16 precision. Entity embeddings are computed using all-MiniLM-L6-v2 (384-dimensional, L2-normalized). The LLM serves three roles: (1) entity extraction from source documents, (2) answer generation given retrieved context, and (3) community summary generation for global QA retrieval.

Methods. We compare six methods in a progressive ablation design:

- **B1 (Rebuild):** Full graph reconstruction from scratch—re-processes all accumulated documents. Quality upper bound, maximum cost.
- **B2 (Naive):** LightRAG-style set-union incremental—processes only new documents, no post-processing.
- **B3 (+Dedup):** B2 + semantic deduplication ($\tau = 0.95$).
- **B4 (+Relink):** B3 + embedding-based graph relinking ($\text{sim} \geq 0.75$).
- **IGV (Ours):** B4 + incremental community repartitioning + community summaries.
- **VectorRAG:** Embedding-based vector retrieval baseline (no graph).

Evaluation. Two modes: (1) *Factoid QA*—100 dataset-specific questions per dataset, evaluated by Recall@5 (R@5) and Answer F1; (2) *Global QA*—20 LLM-generated abstract questions per dataset (e.g., “How many films directed by Tim Burton are mentioned in the corpus?”), evaluated by Answer F1. Global questions are generated by providing 20 sample passages to the LLM and requesting abstract questions requiring cross-document synthesis.

Fair Ablation Design. A critical methodological contribution is that all graph-based methods (B1–IGV) share *cached* LLM entity extractions. Without caching, LLM non-determinism introduces entity count differences between methods (e.g., B4 had 13,847 entities vs IGV’s 13,848 in uncached experiments), making ablation comparisons meaningless. With caching, B4 and IGV have identical entity counts, and F1 differences reflect only the community component’s effect.

Efficiency Metric. We report Update Efficiency Ratio (UER) = B1 total time / method total time, where total time includes both LLM entity extraction and graph construction. For B1, extraction time = $N \times t_{\text{extract}}$ (re-extracting all passages). For incremental methods, extraction time = $|D_t| \times t_{\text{extract}}$ (extracting only the new batch).

4.2 Main Results

Table 2 presents results across all five datasets. Key observations:

IGV matches factoid QA quality. On all datasets, IGV achieves factoid F1 within 0.005 of B2, confirming that incremental updates with deduplication and relinking do not degrade retrieval quality. VectorRAG outperforms all graph methods on factoid QA (F1=0.442–0.570 vs 0.111–0.447), which is expected—vector retrieval excels at factoid lookups.

IGV improves global QA on 3/5 datasets. Community summary retrieval provides substantial improvements on HotpotQA (+12.6%), 2Wiki (+93.8%), and StreamingQA (+38.5%). On StreamingQA, IGV even surpasses VectorRAG (0.281 vs 0.273).

Efficiency is consistent. UER=10.0× across all datasets and methods, as the 70/30 base/batch split yields a consistent 10:1 ratio between full re-extraction and incremental extraction costs.

4.3 Global QA: Community Summary Advantage

Figure 1 shows global QA F1 across datasets. IGV outperforms naive incremental (B2) on 3/5 datasets:

- **HotpotQA:** IGV achieves F1=0.129 vs B2’s 0.114 (+12.6%). Community summaries help answer abstract questions about corpus themes, such as counting entities of a certain type or identifying relationships across documents.
- **2Wiki:** IGV achieves F1=0.375 vs B2’s 0.194 (+93.8%). The largest improvement, demonstrating that community structure is especially valuable when entities have strong topical clustering. 2Wiki’s entity-centric questions benefit from community-level thematic grouping.
- **StreamingQA:** IGV achieves F1=0.281 vs B2’s 0.203 (+38.5%). Notably, IGV also *surpasses VectorRAG* (0.273), showing that graph community structure provides an advantage over flat vector retrieval on streaming corpora where temporal and topical coherence matters.

On MuSiQue and NarrativeQA, IGV underperforms B2. MuSiQue requires precise 2–4 hop compositional reasoning where community summaries are too coarse-grained—they direct retrieval to topically related but reasoning-irrelevant passages. NarrativeQA’s narrative structure may not align well with Louvain’s topological communities, as narrative connections are often temporal rather than entity-based.

4.4 Factoid QA: Quality Preservation

Figure 2 shows factoid QA F1. IGV matches B2 on all datasets (F1 difference <0.005), confirming that incremental updates with deduplication and relinking do not degrade factoid retrieval quality. This is expected: the dual-level retrieval used for factoid QA is identical across B2, B3, B4, and IGV (since all share the same cached entity extraction). The community bonus in IGV’s global retrieval mode does not affect factoid retrieval, which uses a separate code path. VectorRAG outperforms all graph methods on factoid QA, which

Table 2: Main results (5 datasets, 6 methods). Factoid: 100 questions. Global: 20 questions. Best graph method in bold.

Dataset	Method	Factoid R@5	Factoid F1	Global F1	Entities	Communities	UER
HotpotQA	B1 Rebuild	0.485	0.447	0.114	14,721	0	1.0×
	B2 Naive	0.485	0.447	0.114	14,721	0	10.0×
	B3 +Dedup	0.485	0.447	0.114	14,697	0	10.0×
	B4 +Relink	0.485	0.447	0.114	14,697	0	10.0×
	IGV (Ours)	0.485	0.447	0.129	14,697	455	10.0×
	VectorRAG	0.735	0.570	0.154	—	—	—
2Wiki	B1 Rebuild	0.443	0.282	0.194	13,905	0	1.0×
	B2 Naive	0.443	0.282	0.194	13,905	0	10.0×
	B3 +Dedup	0.450	0.272	0.194	13,864	0	10.0×
	B4 +Relink	0.450	0.279	0.194	13,864	0	10.0×
	IGV (Ours)	0.450	0.279	0.375	13,864	612	10.0×
	VectorRAG	0.610	0.361	0.396	—	—	—
MuSiQue	B1 Rebuild	0.350	0.111	0.206	11,795	0	1.0×
	B2 Naive	0.350	0.111	0.206	11,795	0	10.0×
	B3 +Dedup	0.355	0.111	0.234	11,779	0	10.0×
	B4 +Relink	0.355	0.111	0.234	11,779	0	10.0×
	IGV (Ours)	0.355	0.111	0.121	11,779	523	10.0×
	VectorRAG	0.540	0.177	0.292	—	—	—
NarrativeQA	B1 Rebuild	0.475	0.387	0.189	8,956	0	1.0×
	B2 Naive	0.475	0.387	0.189	8,956	0	10.0×
	B3 +Dedup	0.475	0.387	0.189	8,956	0	10.0×
	B4 +Relink	0.475	0.387	0.189	8,956	0	10.0×
	IGV (Ours)	0.475	0.387	0.171	8,956	236	10.0×
	VectorRAG	0.735	0.570	0.154	—	—	—
StreamingQA	B1 Rebuild	0.370	0.315	0.203	15,170	0	1.0×
	B2 Naive	0.370	0.315	0.203	15,170	0	10.0×
	B3 +Dedup	0.370	0.317	0.203	15,132	0	10.0×
	B4 +Relink	0.370	0.320	0.203	15,132	0	10.0×
	IGV (Ours)	0.370	0.320	0.281	15,132	548	10.0×
	VectorRAG	0.595	0.442	0.273	—	—	—




Figure 1: Global QA F1 across datasets. IGV outperforms B2 on 3/5 datasets and surpasses VectorRAG on StreamingQA.




Figure 2: Factoid QA F1. IGV matches B2 (quality preserved) with 10× efficiency over B1.

4.5 Efficiency: 10× Speedup

Figure 3 compares total update time (extraction + construction). B1 rebuild requires re-extracting all passages (~110 min for 2,000 passages at ~3.3s/passage), while IGV extracts only the new batch (~11 min for 200 passages). This yields

is expected—vector retrieval excels at factoid lookups, while graph retrieval’s advantage lies in global reasoning.




Figure 3: Total update time. IGV is 10 $\times$ faster than full rebuild (UER=10.0).

UER=10.0 $\times$ consistently across all datasets. The graph construction overhead (dedup + relink + community detection + summary generation) adds only 0–8 seconds, negligible compared to extraction time. As the corpus grows, B1’s cost grows linearly ($O(N)$), while IGV’s remains proportional to batch size ($O(m)$), making IGV increasingly advantageous for long-running streaming scenarios.

4.6 Ablation Study

Figure 4 shows IGV’s global F1 improvement over B2, isolating each component’s contribution. The progressive ablation B2→B3→B4→IGV reveals:

- **Deduplication (B2→B3):** On MuSiQue, B3 improves global F1 from 0.206 to 0.234 (+13.6%), showing that removing duplicate entities reduces retrieval noise. On other datasets, deduplication has minimal effect on global F1 because the conservative $\tau = 0.95$ merges few entities. Entity count is reduced by 0.1–0.3%.
- **Relinking (B3→B4):** Relinking has minimal effect on global F1 across all datasets. On 2Wiki and StreamingQA, B4 matches B3 exactly. This is because relinking adds edges that improve graph connectivity but do not change community assignments (communities are determined by the full graph topology, not individual edges).
- **Community summaries (B4→IGV):** This is the most impactful component. On 2Wiki, IGV improves global F1 from 0.194 to 0.375 (+93.8%). On StreamingQA, from 0.203 to 0.281 (+38.5%). On HotpotQA, from 0.114 to 0.129 (+12.6%). The community summary bonus provides a thematic signal that factoid retrieval alone cannot capture. On MuSiQue, however, community summaries hurt (−48.3% from B4’s 0.234), confirming that coarse-grained thematic grouping is harmful for precise multi-hop reasoning.




Figure 4: IGV global F1 improvement over B2. Positive on 3/5 datasets.




Figure 5: Entity count (left) and community count (right). Dedup reduces redundancy; communities provide thematic index.

4.7 Entity Deduplication and Community Statistics

Figure 5 shows entity counts and community counts. Deduplication reduces entity count by 0.1–0.3% (conservative $\tau = 0.95$), primarily merging case variants (e.g., “Winner”/“WINNER”) and near-duplicate surface forms. IGV maintains 236–612 communities across datasets, providing the thematic structure for global QA retrieval. The community count varies with corpus size and entity density: HotpotQA (455 communities, 14,697 entities) and StreamingQA (548 communities, 15,132 entities) have similar entity counts but different community structures, reflecting different topical organization.

5 DISCUSSION

When community summaries help. Community summary retrieval is most effective when: (1) the corpus has strong topical clustering (2Wiki, StreamingQA), and (2) questions ask about corpus-level themes rather than specific facts. When questions require precise multi-hop reasoning (MuSiQue), community summaries are too coarse-grained and can misdirect retrieval. This suggests an adaptive approach: use question classification to select between factoid and global retrieval modes, applying community summaries only when the question is thematic.

Factoid vs. global QA trade-off. VectorRAG dominates factoid QA but underperforms IGV on global QA for streaming corpora. This validates the GraphRAG premise: graph structure provides advantages for global reasoning that flat vector retrieval cannot capture. The trade-off suggests that production systems should use hybrid retrieval: vector retrieval for factoid questions and graph community retrieval for global questions.

Fair ablation design. A key methodological contribution is cached entity extraction. Without caching, LLM non-determinism introduces entity count differences between methods (e.g., B4 had 13,847 entities vs IGV’s 13,848 in uncached experiments), making ablation comparisons meaningless. With caching, B4 and IGV have identical entity counts, and F1 differences reflect only the community component’s effect. This insight is broadly applicable to any GraphRAG ablation study using LLM-based extraction.

Limitations. (1) Single random seed; multi-seed evaluation with significance testing would strengthen claims. (2) Global QA questions are LLM-generated and may not cover all question types; human-annotated global QA benchmarks would provide more reliable evaluation. (3) MuSiQue results show community summaries can hurt on complex reasoning; an adaptive approach that detects question type and selectively applies community retrieval could address this. (4) The corpus scale (2,000 passages) is smaller than production deployments; while our complexity analysis shows IGV’s advantages increase with scale, empirical validation at 10K+ passages is needed.

Future work. (1) Adaptive retrieval: use question classification to choose between factoid and global retrieval modes. (2) Multi-seed evaluation with paired t-test and Cohen’s d . (3) Larger-scale experiments (10K+ passages) to demonstrate scalability. (4) Improved community summaries using graph structure (e.g., relation paths) rather than just entity lists. (5) Integration with production GraphRAG systems (LightRAG, GraphRAG) as an incremental maintenance module.

6 CONCLUSION

We presented IGV, an incremental GraphRAG framework that achieves 10× efficiency over full rebuild while maintaining factoid QA quality. Through community summary retrieval, IGV improves global QA F1 by 12.6–93.8% over naive incremental on 3/5 datasets, and surpasses VectorRAG on streaming corpora. A critical methodological insight is that fair ablation requires shared LLM entity extraction caching, eliminating the non-determinism that invalidated prior comparisons. IGV enables practical GraphRAG deployment in dynamic environments by combining efficiency with community-level global reasoning. The trade-off between factoid and global QA performance suggests that hybrid retrieval systems—combining vector retrieval for factoid questions with graph community retrieval for global questions—offer the best of both worlds.

REFERENCES

- [1] Garima Agrawal, Tharindu Kumarage, Zeyad Alghamdi, and Huan Liu. 2024. Can Knowledge Graphs Reduce Hallucinations in LLMs? : A Survey. In *NAACL 2024*. 3947–3960.
- [2] Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, H Niewiadomski, Piotr Nyczyk, and Torsten Hoeffler. 2024. Graph of Thoughts: Solving Elaborate Problems with Large Language Models. In *AAAI 2024*, Vol. 38. 17682–17690.
- [3] Vincent D Blondel, Jean-Loup Guillaume, Renaud Lambiotte, and Etienne Lefebvre. 2008. Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment* 2008, 10 (2008), P10008.
- [4] Chenyang Bu, Guojie Chang, Zihao Chen, CunYuan Dang, Zhize Wu, Yi He, and Xindong Wu. 2025. Query-Driven Multimodal GraphRAG: Dynamic Local Knowledge Graph Construction for Online Reasoning. In *ACL 2025 Findings*. 21360–21380.
- [5] Yukun Cao, Zengyi Gao, Zhiyang Li, Xike Xie, S. Kevin Zhou, and Jianliang Xu. 2024. LEGO-GraphRAG: Modularizing Graph-based Retrieval-Augmented Generation for Design Space Exploration. *Proceedings of the VLDB Endowment* (2024).
- [6] Boyu Chen, Zirui Guo, Zidan Yang, Yuluo Chen, Junze Chen, Zhenghao Liu, Chuan Shi, and Cheng Yang. 2026. PathRAG: Pruning Graph-based Retrieval Augmented Generation with Relational Paths. In *AAAI 2026*, Vol. 40. 30183–30191.
- [7] Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, Dasha Metropolitanaky, Robert Osazuwa Ness, and Jonathan Larson. 2024. From Local to Global: A GraphRAG Approach to Query-Focused Summarization. *arXiv preprint arXiv:2404.16130* (2024).
- [8] Yan Fan, Yu Wang, Pengfei Zhu, and Qinghua Hu. 2024. Dynamic Sub-graph Distillation for Robust Semi-supervised Continual Learning. In *AAAI 2024*, Vol. 38. 11927–11935.
- [9] Yifu Gao, Linbo Qiao, Zhigang Kan, Zhihua Wen, Yongquan He, and Dongsheng Li. 2024. Two-stage Generative Question Answering on Temporal Knowledge Graph Using Large Language Models. In *ACL 2024 Findings*. 6719–6734.
- [10] Zirui Guo, Lianghao Xia, Yanhua Yu, Tu Ao, and Chao Huang. 2025. LightRAG: Simple and Fast Retrieval-Augmented Generation. In *EMNLP 2025 Findings*. 10746–10761.
- [11] Haoyu Han, Yu Wang, Harry Shomer, Kai Guo, Jiayuan Ding, Yongjia Lei, Mahantesh Halappanavar, Ryan A Rossi, Subhabrata Mukherjee, Xianfeng Tang, et al. 2024. Retrieval-Augmented Generation with Graphs (GraphRAG). *arXiv preprint arXiv:2501.00309* (2024).
- [12] Xiaoxin He, Yijun Tian, Yifei Sun, Nitesh V Chawla, Thomas Laurent, Yann LeCun, Xavier Bresson, and Bryan Hooi. 2024. G-Retriever: Retrieval-Augmented Generation for Textual Graph Understanding and Question Answering. In *NeurIPS 2024*.
- [13] Marvin Hofer, Daniel Obraczka, Alieh Saeedi, Hanna Köpcke, and Erhard Rahm. 2024. Construction of Knowledge Graphs: Current State and Challenges. *Information* 15, 8 (2024), 509.
- [14] Yuntong Hu, Zhihan Lei, Zheng Zhang, Bo Pan, Chen Ling, and Liang Zhao. 2025. GRAG: Graph Retrieval-Augmented Generation. In *NAACL 2025 Findings*. 4145–4157.
- [15] Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. 2024. HippoRAG: Neurobiologically Inspired Long-Term Memory for Large Language Models. In *NeurIPS 2024*.
- [16] S H Kim, Seok Jun Hwang, JungHyoun Kim, Jeong Hyeon Park, and Yong Suk Choi. 2025. ReGraphRAG: Reorganizing Fragmented Knowledge Graphs for Multi-Perspective Retrieval-Augmented Generation. In *EMNLP 2025 Findings*. 5426–5443.
- [17] Junlin Lee, Yequan Wang, Jing Li, and Min Zhang. 2024. Multimodal Reasoning with Multimodal Knowledge Graph. In *ACL 2024*. 10767–10782.
- [18] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *NeurIPS 2020*.
- [19] Mufei Li, Siqi Miao, and Pan Li. 2025. Simple is Effective: The Roles of Graphs and Large Language Models in Knowledge-Graph-Based Retrieval-Augmented Generation. In *ICLR 2025*.
- [20] Lei Liang, Meng Sun, et al. 2024. KAG: Boosting LLMs in Professional Domains via Knowledge Augmented Generation. *arXiv preprint arXiv:2409.13731* (2024).

- [21] Wanying Liang, Pasquale De Meo, Yong Tang, and Jia Zhu. 2024. A Survey of Multi-modal Knowledge Graphs: Technologies and Trends. *Comput. Surveys* 56, 11 (2024), 1–41.
- [22] Siru Liu, Aileen P Wright, Allison B McCoy, Sean S Huang, Bryan D Steitz, and Adam Wright. 2025. Detecting emergencies in patient portal messages using large language models and knowledge graph-based retrieval-augmented generation. *J. Am. Medical Informatics Assoc.* 32, 6, 1032–1039.
- [23] Antonio Longa, Steve Azzolin, Gabriele Santin, Giulia Cencetti, Pietro Lió, Bruno Lepri, and Andrea Passerini. 2024. Explaining the Explainers in Graph Neural Networks: a Comparative Study. *Comput. Surveys* 57, 5 (2024), 1–37.
- [24] Linhao Luo, Zicheng Zhao, Gholamreza Haffari, Dinh Phung, Chen Gong, and Shirui Pan. 2025. GFM-RAG: Graph Foundation Model for Retrieval Augmented Generation. *arXiv preprint arXiv:2502.01113*.
- [25] Shengjie Ma, Chengjin Xu, Xuhui Jiang, Muzhi Li, Huaren Qu, Cehao Yang, Jiaxin Mao, and Jian Guo. 2025. Think-on-Graph 2.0: Deep and Faithful Large Language Model Reasoning with Knowledge-guided Retrieval Augmented Generation. In *ICLR 2025*.
- [26] Costas Mavromatis and George Karypis. 2025. GNN-RAG: Graph Neural Retrieval for Large Language Model Reasoning on Knowledge Graphs. In *ACL 2025 Findings*. 16682–16699.
- [27] Congmin Min, Sahil Bansal, Joyce Pan, Abbas Keshavarzi, Rhea Mathew, and Amar Viswanathan Kannan. 2025. Towards Practical GraphRAG: Efficient Knowledge Graph Construction and Hybrid Retrieval at Scale. *arXiv preprint arXiv:2507.03226* (2025).
- [28] Boci Peng, Yun Zhu, Yongchao Liu, Xiaohe Bo, Haizhou Shi, Chuntao Hong, Yan Zhang, and Siliang Tang. 2025. Graph Retrieval-Augmented Generation: A Survey. *ACM Transactions on Information Systems* 44, 2 (2025), 1–52.
- [29] Deeksha Prahlad, Chanhee Lee, Dongha Kim, and Hokeun Kim. 2025. Personalizing Large Language Models using Retrieval Augmented Generation and Knowledge Graph. In *WWW 2025*.
- [30] Xinying Qian, Ying Zhang, Yu Zhao, Baohang Zhou, Xuhui Sui, Li Zhang, and Kehui Song. 2024. TimeR4: Time-aware Retrieval-Augmented Large Language Models for Temporal Knowledge Graph Question Answering. In *EMNLP 2024*. 6942–6952.
- [31] Bhaskarjit Sarmah, Dhagash Mehta, Benika Hall, Rohan Rao, Sunil Patel, and Stefano Pasquali. 2024. HybridRAG: Integrating Knowledge Graphs and Vector Retrieval Augmented Generation for Efficient Information Extraction. In *CIKM 2024*. 608–616.
- [32] Jiashuo Sun, Chengjin Xu, Lumingyuan Tang, Saizhuo Wang, Chen Lin, Yeyun Gong, Heung-Yeung Shum, and Jian Guo. 2024. Think-on-Graph: Deep and Responsible Reasoning of Large Language Model on Knowledge Graph. In *ICLR 2024*.
- [33] Xingyu Tan, Xiaoyang Wang, Qing Liu, Xiwei Xu, Xin Yuan, and Wenjie Zhang. 2025. Paths-over-Graph: Knowledge Graph Empowered Large Language Model Reasoning. In *WWW 2025*. 3505–3522.
- [34] Shiyu Tian, Shuyue Xing, Xingrui Li, Yangyang Luo, Caixia Yuan, Wei Chen, Huixing Jiang, and Xiaojie Wang. 2025. A Systematic Exploration of Knowledge Graph Alignment with Large Language Models in Retrieval Augmented Generation. In *AAAI 2025*.
- [35] Yijun Tian, Huan Song, Zichen Wang, Haozhu Wang, Ziqing Hu, Fang Wang, Nitesh V Chawla, and Panpan Xu. 2024. Graph Neural Prompting with Large Language Models. In *AAAI 2024*, Vol. 38. 19080–19088.
- [36] Vincent A Traag, Ludo Waltman, and Nees Jan van Eck. 2019. From Louvain to Leiden: guaranteeing well-connected communities. *Scientific Reports* 9 (2019), 5233.
- [37] Shijie Wang, Wenqi Fan, Yue Feng, Shanru Lin, Xinyu Ma, Shuaiqiang Wang, and Dawei Yin. 2025. Knowledge Graph Retrieval-Augmented Generation for LLM-based Recommendation. In *ACL 2025*. 27152–27168.
- [38] Yibo Wang, Haoyang Li, Fei Teng, and Lei Chen. 2025. GORAG: Graph-based Online Retrieval Augmented Generation for Dynamic Few-shot Social Media Text Classification. *arXiv preprint arXiv:2501.02844* (2025).
- [39] Zhizheng Wang, Yuanyuan Sun, Xiaokun Zhang, Bo Xu, Zhihao Yang, and Hongfei Lin. 2024. Continual learning with high-order experience replay for dynamic network embedding. *Pattern Recognition* 159 (2024), 111093.
- [40] Yilin Wen, Zifeng Wang, and Jimeng Sun. 2024. MindMap: Knowledge Graph Prompting Sparks Graph of Thoughts in Large Language Models. In *ACL 2024*. 10370–10388.
- [41] Junde Wu, Jiayuan Zhu, Qi Yang, Jingkun Chen, Min Xu, Filippo Menolascina, and Vicente Grau. 2025. Medical Graph RAG: Towards Safe Medical Large Language Model via Graph Retrieval-Augmented Generation. In *ACL 2025*.
- [42] Tongtong Wu, Linhao Luo, Yuan-Fang Li, Shirui Pan, Thuy-Trang Vu, and Gholamreza Haffari. 2024. Continual Learning for Large Language Models: A Survey. *arXiv preprint arXiv:2402.01364* (2024).
- [43] Zhentao Xu, Missy Meine C Dela Cruz, Matthew Guevara, Tie Wang, M N Deshpande, Xiaofeng Wang, and Zheng Li. 2024. Retrieval-Augmented Generation with Knowledge Graphs for Customer Service Question Answering. In *KDD 2024*. 2905–2909.
- [44] Xu Yao, Shizhu He, Jiabei Chen, Zihao Wang, Yangqiu Song, Hanghang Tong, Guang Liu, Jun Zhao, and Kang Liu. 2024. Generate-on-Graph: Treat LLM as both Agent and KG for Incomplete Knowledge Graph Question Answering. In *EMNLP 2024*. 18410–18430.
- [45] Qiao Yuan, Sheng-Wei Guan, Pin Ni, Tianlun Luo, Prudence Wong, Victor Chang, and Ka Lok Man. 2026. Continual graph learning: A survey. *Pattern Recognition* 179 (2026), 113600.
- [46] Bowen Zhang and Harold Soh. 2024. Extract, Define, Canonicalize: An LLM-based Framework for Knowledge Graph Construction. In *EMNLP 2024*. 9820–9836.
- [47] Xinni Zhang, Yankai Chen, Chenhao Ma, Yixiang Fang, and Irwin King. 2024. Influential Exemplar Replay for Incremental Learning in Recommender Systems. In *AAAI 2024*, Vol. 38. 9368–9376.
- [48] Xuejiao Zhao, Siyan Liu, Su-Yin Yang, and Chunyan Miao. 2025. MedRAG: Enhancing Retrieval-augmented Generation with Knowledge Graph-Elicited Reasoning for Healthcare Copilot. In *WWW 2025*. 4442–4457.
- [49] Yanping Zheng, Yi Lü, and Zhewei Wei. 2024. A survey of dynamic graph neural networks. *Frontiers of Computer Science* 19, 6 (2024).
- [50] Xiangrong Zhu, Yuexiang Xie, Yi Liu, Yuanman Li, and Wei Hu. 2025. Knowledge Graph-Guided Retrieval Augmented Generation. In *NAACL 2025*. 8912–8924.